

License Plate Recognition

阮柏愷 bkruan.ee11@nycu.edu.tw

黃柏綸 kevin503.ee12@nycu.edu.tw

Special Thanks to Yi-Lun Wu 🙏

Outline

- Problem Description
- Dataset
- Tasks
 - Setting up Environment
 - Build Dataset
 - Define Model
 - Save Model

Problem Description

- Input: An image of a License plate.
- Output: A string of length 7.
- We can make some assumptions:
 - All license plates are centered in the images.
 - The number of characters in a license plate is fixed to 7.
 - Only characters 0~9 and A~Z are legal.



Dataset

- Our dataset was taken from this [work](#) with a few modifications.
- Training set:
 - It contains 8 different folders with a total of 8762 images.
 - `trainValNew.csv` contains all the image paths of the training set.
- Testing set
 - It contains 4 different folders with a total of 299 images.
 - `trainValNew.csv` contains all the image paths of the testing set.
- The CSV format is as follows:

```
track_id , image_path           , lp      , train
s01_l01_1, s01_l01_unique/1_1.png, 7C24698, 1
s01_l01_3, s01_l01_unique/3_1.png, 5AR2347, 1
...
```

Source Code

- SSH into an arbitrarily server.
- Clone the source code from Github. Note, **if the source code has been cloned in the previous lab, there is no need to clone it again this time.**

```
$ git clone https://github.com/Justin900429/hcc-m1
```

- Make sure that your working directory is hcc-m1/lab2

```
$ cd hcc-m1/lab2
```

Or you can open the folder hcc-m1/lab2 in VS Code.

Create Virtual Environment

- Create virtual environment for Lab2.

```
$ python3 -m venv {name_of_your_venv}
```

For example:

```
$ python3 -m venv lab2env
```

How to Use?

- Activate your virtual environment.

```
$ source [name_of_your_venv]/bin/activate
```

For example:

```
$ source lab2env/bin/activate  
(lab2env) $
```

Once activated, the environment name will be displayed in your terminal prompt.

- Install all python packages.

```
(lab2env) $ pip install -r requirements.txt
```

- To use Nvidia GPU in PyTorch, you need to install the Nvidia driver and the CUDA toolkit. These packages have already been installed on all servers.

Data Preprocessing

- Download the dataset from the Google Drive. [link](#)
Alternatively, you can run `download.py` script to download the dataset from the terminal.

```
$ python download.py
```

- Unzip the dataset.

```
$ unzip dataset.zip
```


Data Preprocessing

- The structure of the ./dataset directory should look like this:

```
dataset
├── train
│   ├── s01_l01_unique
│   ├── s01_l02_unique
│   ├── ...
│   └── trainValNew.csv
└── val
    ├── crop_m1_unique
    ├── crop_m2_unique
    ├── ...
    └── trainValNew.csv
```

Customize dataset in PyTorch

- The base dataset is defined in ./dataset.py.

```
class LicensePlateDataset(torch.utils.data.Dataset):
    def __init__(self, csv_path, transform=None):
        self.transform = transform
        csv = pd.read_csv(csv_path)
        root = os.path.dirname(csv_path)
        self.paths = csv['image_path'].apply(
            lambda p: os.path.join(root, p)).values
        self.labels = csv['lp'].apply(
            lambda lp: [CODES_TO_INT[c] for c in lp]).values

    def __len__(self):
        ...

    def __getitem__(self, idx):
        ...
```

Customize dataset in PyTorch

- The base dataset is defined in `./dataset.py`.

```
class LicensePlateDataset(torch.utils.data.Dataset):
    def __init__(self, csv_path, transform=None):
        ...

    def __len__(self):
        # return the size of the dataset
        return len(self.paths)

    def __getitem__(self, idx):
        # return i-th element in the dataset
        image = Image.open(self.paths[idx]) # load on demand
        if self.transform is not None:
            image = self.transform(image)
        return image, torch.tensor(self.labels[idx])
```

Training Set

- Derive a training set from the LicensePlateDataset:

```
class TrainDataset(LicensePlateDataset):
    def __init__(self,
                  csv_path='./dataset/train/trainValNew.csv'):
        transform = transforms.Compose([
            transforms.ColorJitter(0.2, 0.2, 0.2, 0.2),
            transforms.RandomRotation(3),
            transforms.RandomResizedCrop(
                (HEIGHT, WIDTH),
                scale=(0.7, 1.0),
                ratio=[2., 2.8]),
            transforms.ToTensor()
        ])
        super().__init__(csv_path, transform)
```

Checkpoint 1 - Testing Set

- To complete the testing set, you need to add a resize transformation to the transform object in `./dataset.py`

```
'''  
Checkpoint 1: Resize image to (HEIGHT, WIDTH)  
Hint: Add resize transformation to the beginning of transform  
'''  
  
transform = transforms.Compose([  
    ???  
    transforms.ToTensor()  
)
```

- HEIGHT and WIDTH are global variables in `dataset.py`, so you don't need to define them again.
- You can find the Resize operation in the [documentation](#).

Checkpoint 1

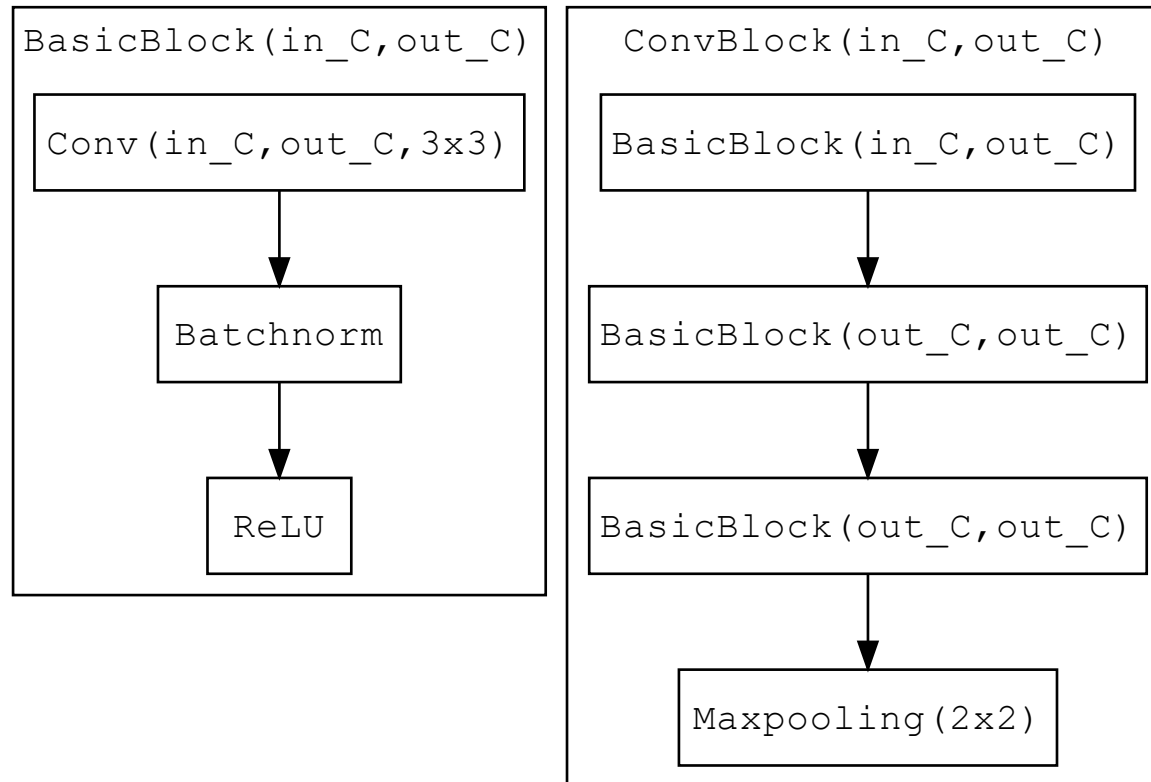
- Execute `dataset.py` to test your modification.

```
$ python dataset.py
```

- **Show** the resolution of the dumped image `val.jpg`.

```
$ file val.jpg
```

Schematic of 2 Basic Modules



Step 1. Define BasicBlock in PyTorch

- Here is a standard implementation of BasicBlock in PyTorch:

```
class BasicBlock(nn.Module):
    def __init__(self, in_channels, out_channels):
        super().__init__()
        self.conv = nn.Conv2d(
            in_channels, out_channels, kernel_size=3, padding=1)
        self.bn = nn.BatchNorm2d(out_channels)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.conv(x)
        x = self.bn(x)
        x = self.relu(x)
        return x
```


Step 1. Alternative Approach for BasicBlock

- Alternatively, you can rewrite BasicBlock using nn.Sequential like this:

```
class BasicBlock(nn.Module):
    def __init__(self, in_channels, out_channels):
        super().__init__()
        self.main = nn.Sequential(
            nn.Conv2d(
                in_channels, out_channels, kernel_size=3, padding=1),
            nn.BatchNorm2d(out_channels),
            nn.ReLU())

    def forward(self, x):
        return self.main(x)
```

Step 2. Define ConvBlock in PyTorch

- Here is a standard implementation of ConvBlock in PyTorch:

```
class ConvBlock(nn.Module):
    def __init__(self, in_channels, out_channels):
        super().__init__()
        self.conv1 = BasicBlock(in_channels, out_channels)
        self.conv2 = BasicBlock(out_channels, out_channels)
        self.conv3 = BasicBlock(out_channels, out_channels)
        self.maxpool = nn.MaxPool2d((2, 2))

    def forward(self, x):
        x = self.conv1(x)
        x = self.conv2(x)
        x = self.conv3(x)
        x = self.maxpool(x)
        return x
```

Checkpoint 2

- You have to rewrite ConvBlock using nn.Sequential.

```
class ConvBlock(nn.Module):  
    '''  
    Checkpoint 2:  
    Use `nn.Sequential` to rewrite `ConvBlock`.  
    '''  
    def __init__(self, in_channels, out_channels):  
        super().__init__()  
        self.main = ???  
  
    def forward(self, x):  
        return ???
```

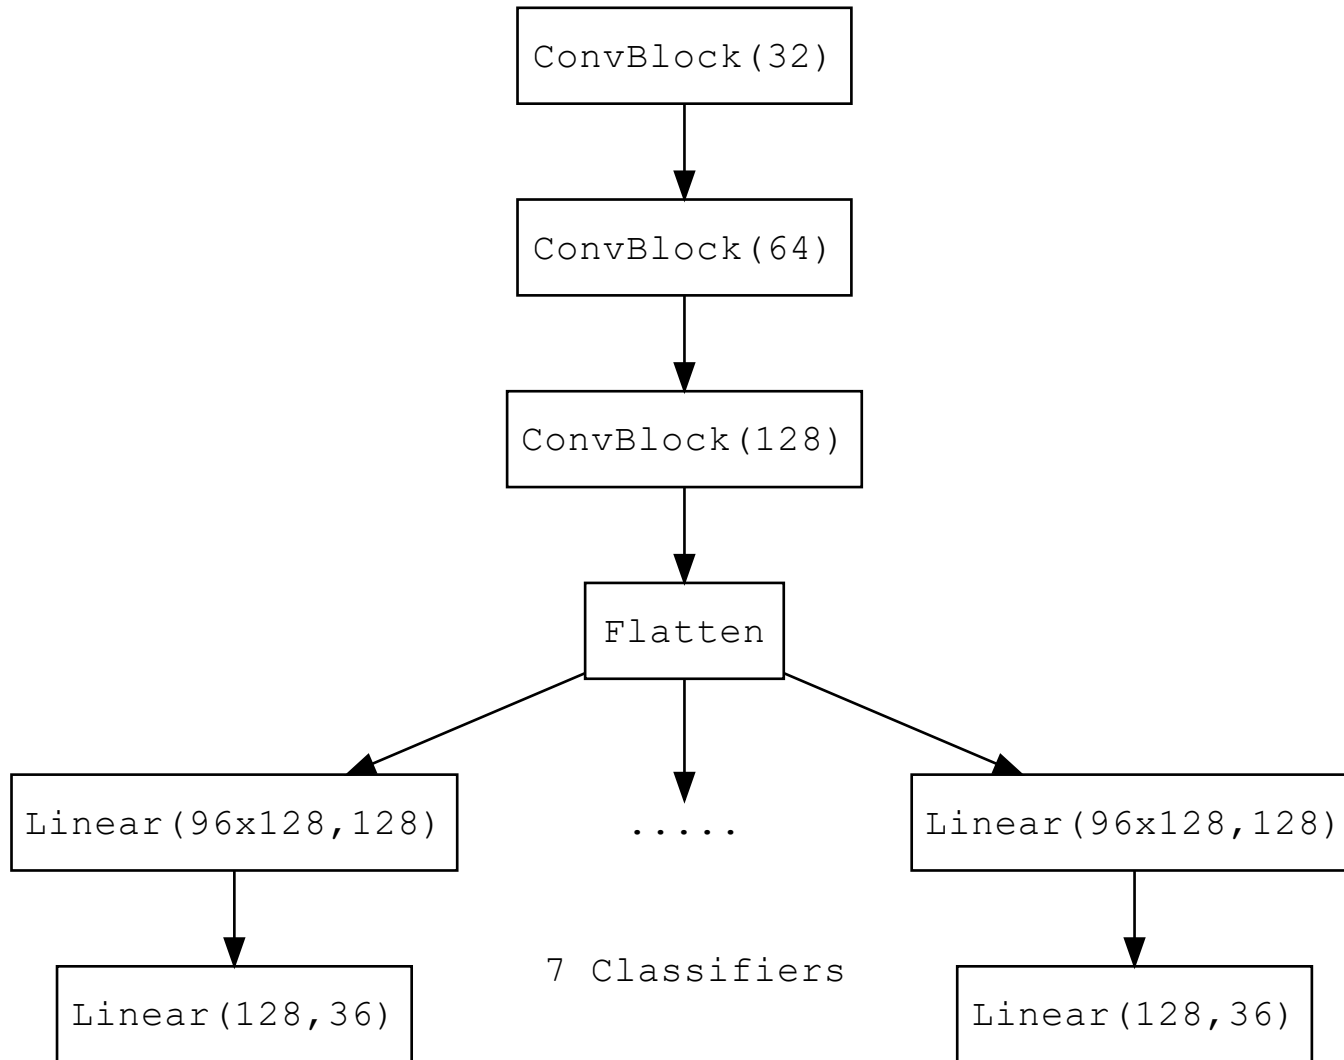
Checkpoint 2

- To test the modules, please run the following command in the terminal after removing the old ConvBlock:

```
$ python model.py
```

- **Show** the output in the terminal.

Schematic of Full Model



Define Full Model

```
class LPRModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.encoder = nn.Sequential(
            ConvBlock(3, 32),          # [B, 32, H / 2, W / 2]
            ConvBlock(32, 64),         # [B, 64, H / 4, W / 4]
            ConvBlock(64, 128),        # [B, 128, H / 8, W / 8]
            nn.Flatten(start_dim=1))   # [B, 128 x H / 8 x W / 8]
        self.classifiers = nn.ModuleList()
        for _ in range(7):
            classifier = nn.Sequential(
                nn.Linear(HEIGHT // 8 * WIDTH // 8 * 128, 128),
                nn.ReLU(),
                nn.Linear(128, 36))
            self.classifiers.append(classifier)
        ...
```

Checkpoint 3

Please complete the forward function of the full model as follows:

```
class LPRModel(nn.Module):
    ...
    def forward(self, x):
        x = self.encoder(x)
        outputs = []
        for classifier in self.classifiers:
            output = classifier(x)          # [B, 36]
            output = output.unsqueeze(1)    # [B, 1, 36]
            outputs.append(output)
        ...
        Checkpoint 3:
            Concatenate `outputs` along the axis 1.
            The shape of the result must be [B, 7, 36].
        ...
        # output = torch.cat(????)
        return output
```

Checkpoint 3

- Download `pretrain.zip` from Google Drive. You can find the link [here](#).

If you have already used `download.py` to download the dataset, you will not need to download `pretrain.zip` as it has already been downloaded.

- To unzip `pretrain.zip`, use the following command:

```
$ unzip pretrain.zip
```

- Then, execute the following command:

```
$ python detect.py \  
  --checkpoint ./pretrain/ckpt100.pt \  
  --image ./example/9B52145.png
```

- Finally, **show** the predictions by showing the following message:

```
Loading image from ./example/9B52145.png  
weight has been loaded  
predictions: 9B52145
```


Training Loop

```
1 Declare Dataset
2 Declare Model
3 Declare Optimizer
4 For epoch in 1..epochs:
5     For batch in Dataset:
6         Calculate loss by forward propagate
7         Calculate gradients by backward propagation
8         Update model
9     end For
10    Evaluate Model
11    Save the checkpoint
12 end For
```

Training Implementation in PyTorch

```
train_loader = DataLoader(dataset=TrainDataset() ...)
val_loader = DataLoader(dataset=ValDataset() ...)
model = LPRModel()
optim = torch.optim.Adam(model.parameters() ...)
...
for epoch in range(args.epochs):
    with tqdm(total=len(train_loader)) as pbar:
        for image, label in train_loader:
            predict = model(image)
            loss = loss_fn(predict.transpose(1, 2), label)
            optim.step()
        ...
    with torch.no_grad():
        for image, label in val_loader:
            ...
    if epoch == 0 or (epoch + 1) % args.checkpoint_period == 0:
        torch.save(model.state_dict(), path)
```

Checkpoint 4

- To execute the training script, run the following command:

```
$ python train.py
```

- As the training process takes around half an hour to complete, it is sufficient to **display** the training progress only. Here's an example of what it would look like:

```
Epoch 1/100 train_loss: 2.3468, val_loss: 2.1914, train_acc: ...  
Epoch 2/100 train_loss: 1.8277, val_loss: 1.8763, train_acc: ...  
Epoch 3/100 train_loss: 1.2882, val_loss: 1.4427, train_acc: ...  
Epoch 4/100 train_loss: 0.9773, val_loss: 1.4076, train_acc: ...  
Epoch 5/100 train_loss: 0.7762, val_loss: 1.0444, train_acc: ...  
...
```

The End