

Object Detection - YOLOv7

阮柏愷 bkruan.ee11@nycu.edu.tw

黃柏綸 kevin503.ee12@nycu.edu.tw

Special Thanks to Yi-Lun Wu 

Outline

- Traditional Object Detection
- R-CNN, Fast R-CNN, Faster-RCNN
- YOLOv7
- Tasks
 - Create a Virtual Environment
 - Inference
 - Annotation
 - Training

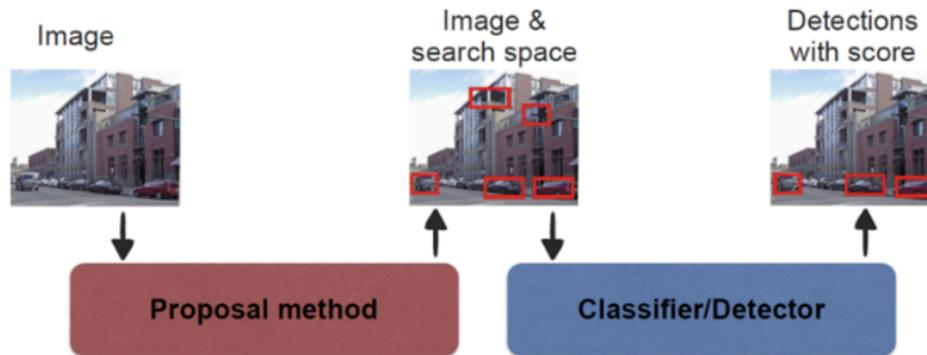
Object Detection

The algorithm produces a list of object categories present in the image, along with an axis-aligned bounding box that indicates the position and scale of every instance.



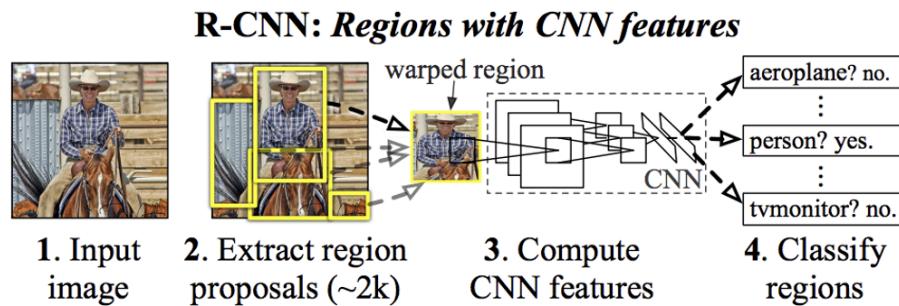
Traditional Approach

1. Object Localization: Hypothesize bounding boxes, e.g., sliding window.
2. Feature Extraction: Extract features for each box, e.g., SIFT, HOG.
3. Image Classification: Classify each box into an object category, e.g., SVM, KNN.



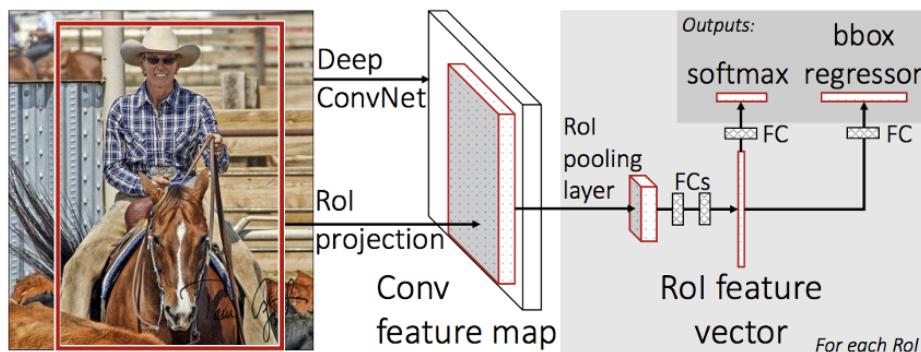
R-CNN-Based Methods

- Use Selective Search to find region proposals.
- Extract features of each proposal using Convolutional Neural Networks.
- Learn the bounding box through regression and predict the object class with SVM.
- Other variants include Fast R-CNN and Faster-RCNN.



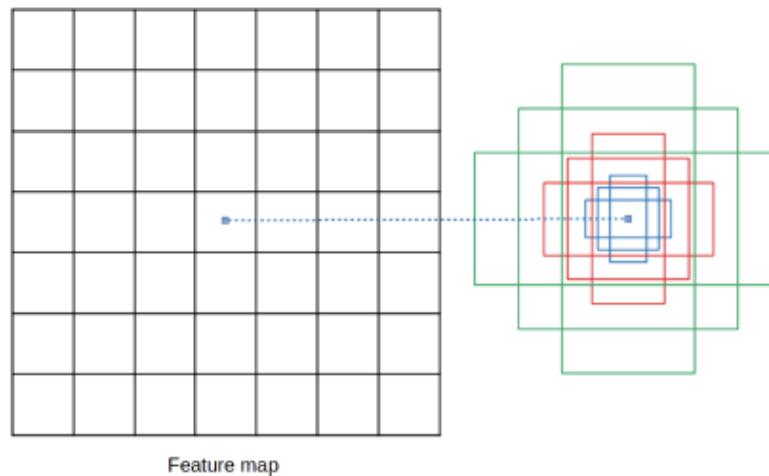
Fast R-CNN:

- Problem with R-CNN
 - Using a deep CNN to extract all region proposals of an image is very slow.
- Solution:
 - Run the CNN only once to get the features of the full image. Features of region proposals are cropped from the overall features.
 - ROI pooling unifies the dimensions of the cropped features.
 - Replace SVM with Neural Networks.



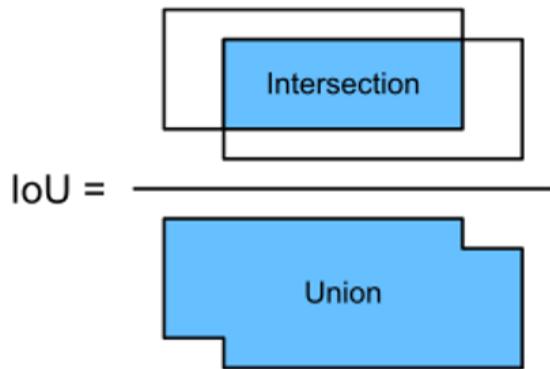
Faster R-CNN

- Problem with Fast R-CNN:
 - Region proposal is time-consuming.
 - Accuracy is highly affected by region proposals.
- Solution:
 - Learn a region proposal algorithm.
 - Introduce anchor boxes to simplify region proposal.



Faster R-CNN

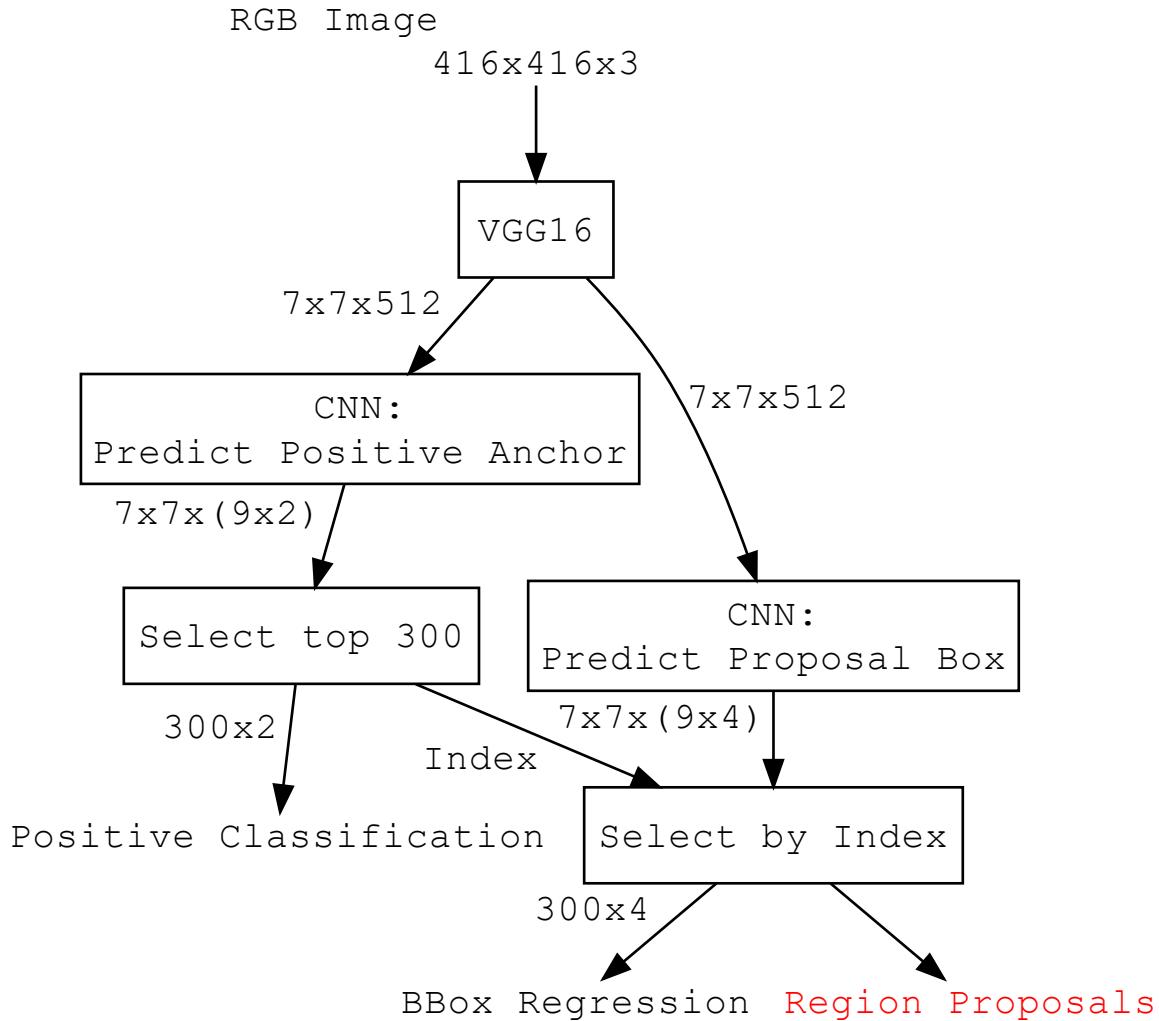
- Intersection over Union (IoU) determines the overlap between a predicted box and a ground-truth box.



- A positive anchor is either:
 1. The anchor with the highest IoU overlap with a ground-truth box.
 2. An anchor that has an IoU overlap higher than 0.7 with any ground-truth box.

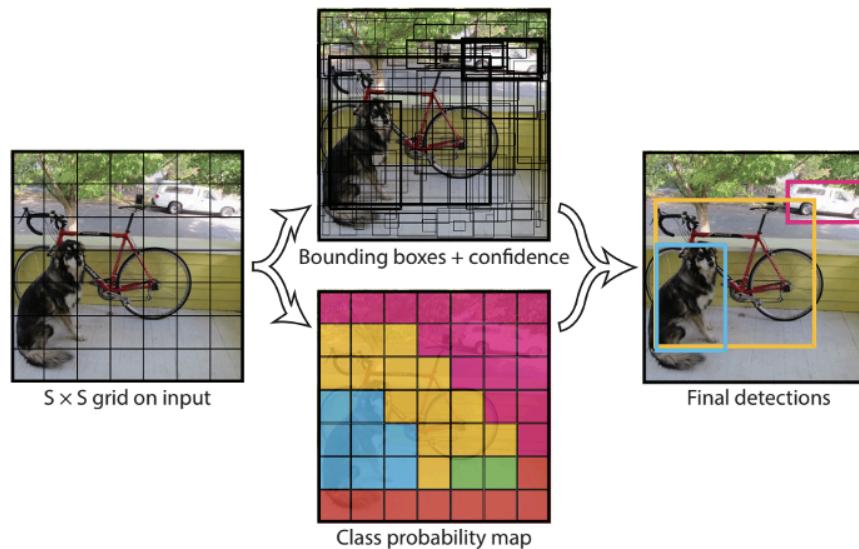
Faster R-CNN

Learn a region proposal model:

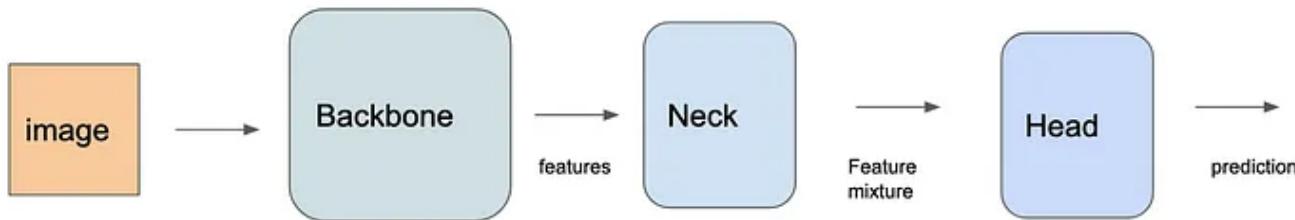


YOLO

- YOLO Family
 - This includes YOLOv1~8, YOLO-NAS, and others.
 - A single neural network is trained end-to-end to predict bounding boxes and class probabilities.



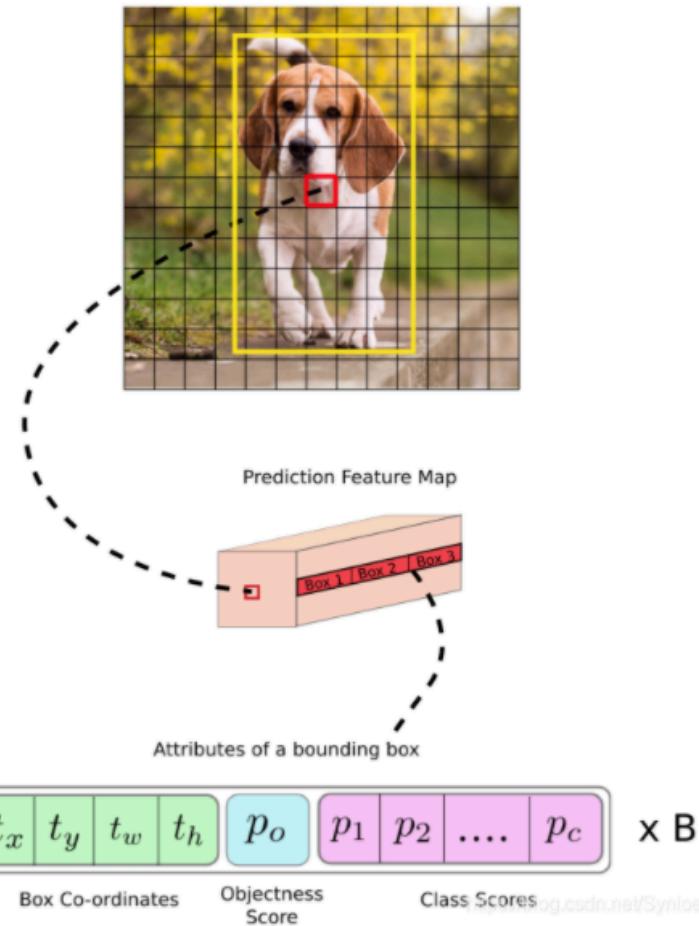
A general YOLO architecture



<https://medium.com/@nahidalam/understanding-yolov7-neural-network-343889e32e4e>

- **Backbone:** Feature Extractor.
- **Neck:** Intermediate layers for fusing features from different level or layers,
- **Head:** Predicting the final results (class and bounding boxes),

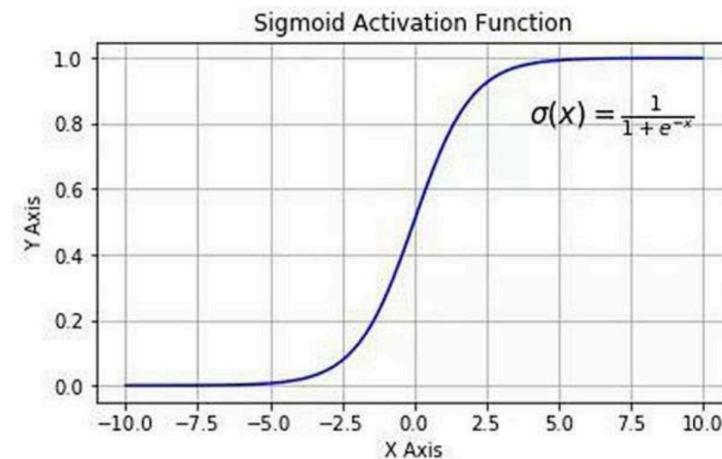
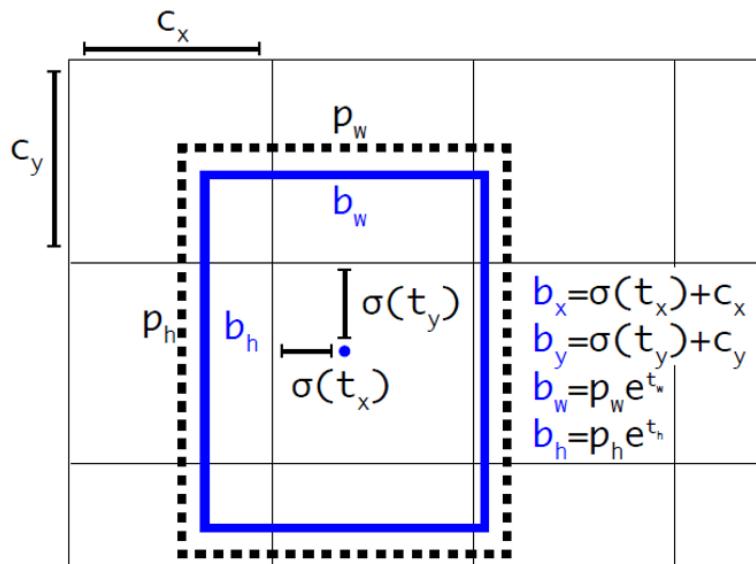
Each Anchor Box Predicts $5 + n_c$ Values



<https://b10515007.medium.com/yolov3-%E8%A9%B3%E7%B4%B0%E8%A7%A3%E8%AA%AA-22e3da36260a>

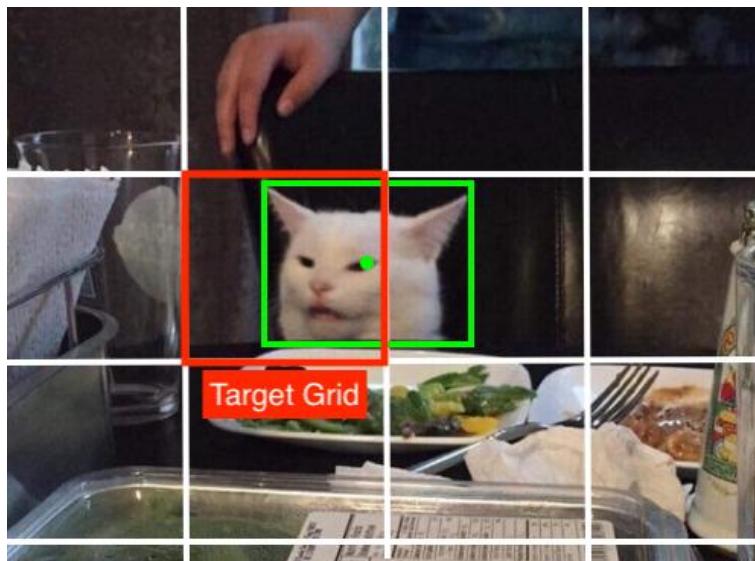
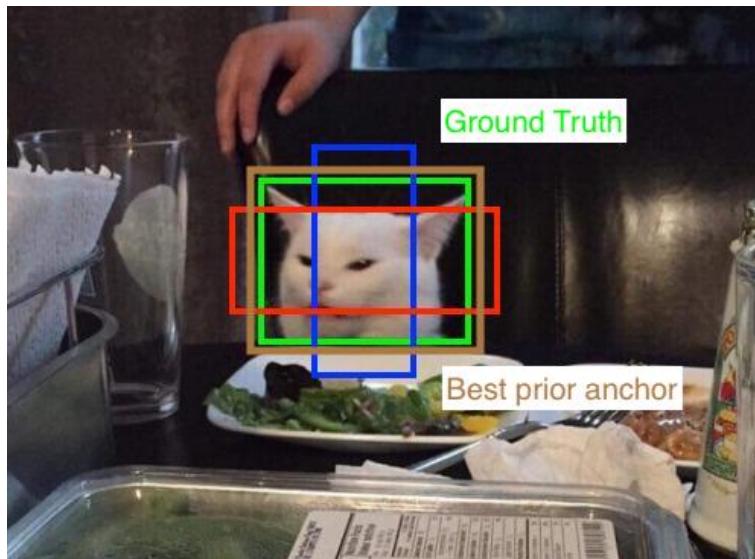
Bounding Box

- Predict the distance from the upper-left corner of the grid to the center of the ground truth bounding box. The predicted value is in the range of $[0, 1]$.
- The size of the bounding box is rescaled from the size of the corresponding prior anchor.



Bounding Box Loss and Class Loss

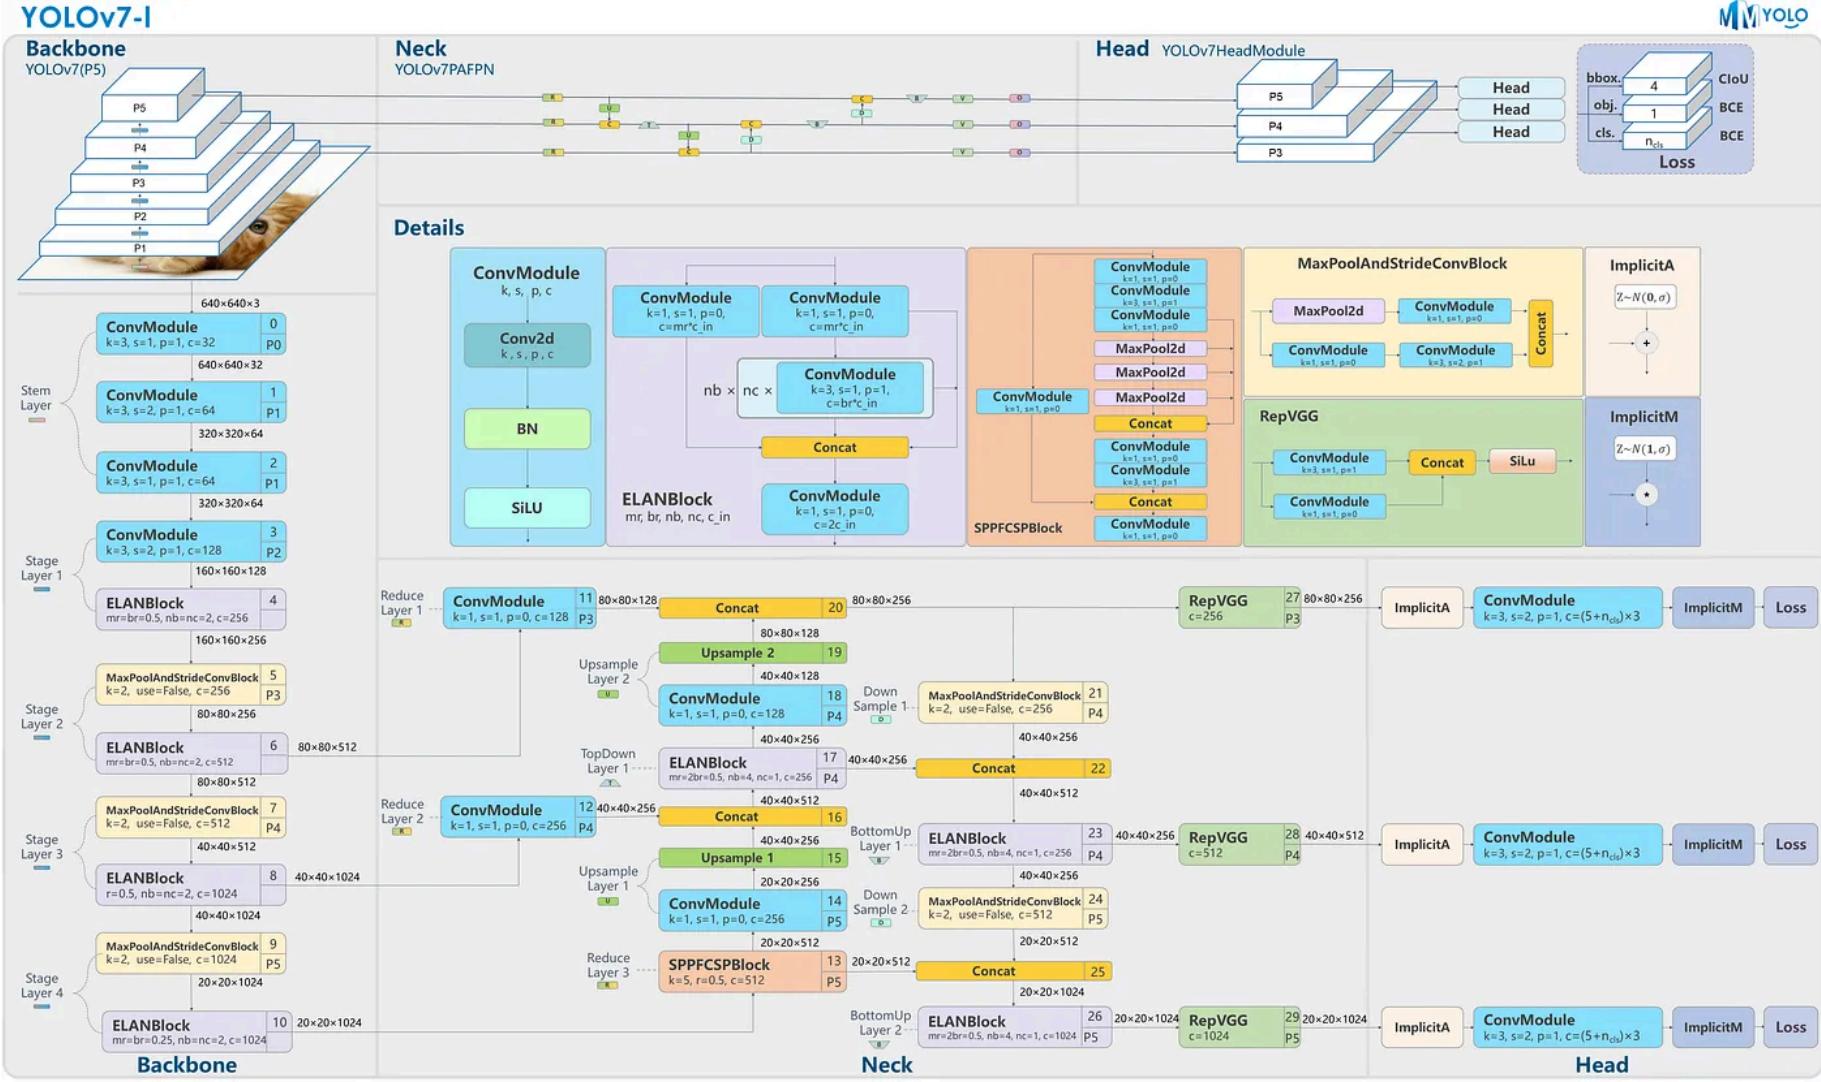
- For each ground truth box, center prior anchors and ground truth then select the one with maximum IoU.
- The selected anchor at the grid where the center of the ground truth locates is called the positive anchor, and the label of this anchor is the corresponding ground truth.
- The other anchors (also called negative anchors) do not contribute to the loss calculation in this step.



Objectness Score (Probability p_o)

- $p_o = 1$ if it is a positive anchor of some ground truth.
- $p_o = 0$ if the maximum IoU between the predicted bounding box and all ground truth is less than 0.5.
- The other anchors do not contribute to the loss calculation in this step.

YOLOv7 Model Architecture



<https://github.com/open-mmlab/mmyolo>

What's new in YOLOv7?

1. **Extended-Efficient Layer Aggregation Network (only in YOLOv7-E6E)**
→ Solving the deteriorate converge when scaling a model.
2. **Compound model scaling**
→ Using depth and width scaling in different model blocks.
3. **Bag of Freebies (BoF)**
→ BoF refers to the techniques that improve the model without increasing training cost. These techniques include Re-parameterization Convolution, Label Assignment with Varying granularity
4. **and more**

Loss function in YOLOv7

- Also address the class and the bounding box with some designs derived from: YOLOv4, YOLOv5, SimOTA.
- The prediction of each anchor includes: (x, y, w, h, objectness score, number of classes)

Source Code

- SSH into an arbitrarily server.
- Clone the source code from Github if you have not already done so:

```
$ git clone https://github.com/Justin900429/hcc-ml
```

Otherwise, pull the updated code from Github using `git pull`

```
$ cd hcc-ml  
$ git pull
```

- Make sure that your working directory is `lab3`.

```
$ cd lab3
```

Create Virtual Environment

- Create a virtual environment for Lab3.

```
$ python3.12 -m venv {name_of_your_venv}
```

For example:

```
$ python3.12 -m venv lab3env
```

- Activate your virtual environment.

```
$ source {name_of_your_venv}/bin/activate
# For example
$ source lab3env/bin/activate
```

- Update pip and install all required Python packages:

```
(lab3env) $ pip install --upgrade pip
(lab3env) $ pip install -r requirements.txt
```

Download Pretrained Models

- Run the following commands to download all pretrained weights:

```
$ cd weights  
$ ./download_weights.sh  
$ cd ..
```

Checkpoint 1

- Run demo.py to test the environment and the pretrained weight.

```
$ python demo.py \
--image ./data/street.jpg \
--weights ./weights/yolov7.pt
```

- Show demo.png



Checkpoint 2

- Use the flag `--help` to see more options of `demo.py`:

```
$ python demo.py --help
usage: demo.py [-h] [--image IMAGE] [--weights WEIGHTS [WEIGHTS ...]
               [--img_size IMG_SIZE] [--conf_threshold CONF_THRESHOLD]
               [--nms_threshold NMS_THRESHOLD] [--device DEVICE]

options:
  -h, --help            show this help message and exit
  --image IMAGE         source
  --weights WEIGHTS [WEIGHTS ...]
                        model.pt path(s)
  --img_size IMG_SIZE   inference size (pixels)
  --conf_threshold CONF_THRESHOLD
                        object confidence threshold
  --nms_threshold NMS_THRESHOLD
                        IOU threshold for NMS
  --device DEVICE       cuda device, i.e. 0 or 0,1,2,3 or cpu
```

Checkpoint 2

- Confidence (`--conf_threshold`):

The confidence of a bbox is the objectness score \times class probability. This is used to filter out the bboxes before Non-Maximum Suppression.

- Non-Maximum Suppression (`--nms_threshold`):

1. Let A be the set of all candidate bboxes and $B = \emptyset$ be the initial result.
2. Repeat steps 3 and 4 until A is empty.
3. Select bbox b in A with the highest confidence. If the maximum IoU between b and all bboxes in B is less than `nms_threshold`, add b to set B .
4. Remove b from A .

Checkpoint 2-1

- Change the confidence threshold to detect the traffic light on the right-hand side of the image.

```
$ python demo.py \
  --image ./data/street.jpg \
  --weights ./weights/yolov7.pt \
  --conf_threshold ??
```

Checkpoint 2-1



Checkpoint 2-2

- Change the confidence threshold and NMS threshold to detect the motorcycles on the rightmost side of the image (there should be at least three motorcycles).

```
$ python demo.py \
--image ./data/street.jpg \
--weights ./weights/yolov7.pt \
--conf_threshold ?? \
--nms_threshold ??
```

Checkpoint 2-2



Label Studio

 Stars 22k

Step 1

The following steps should be executed on your personal computer:

- Install Label Studio by running the following commands:

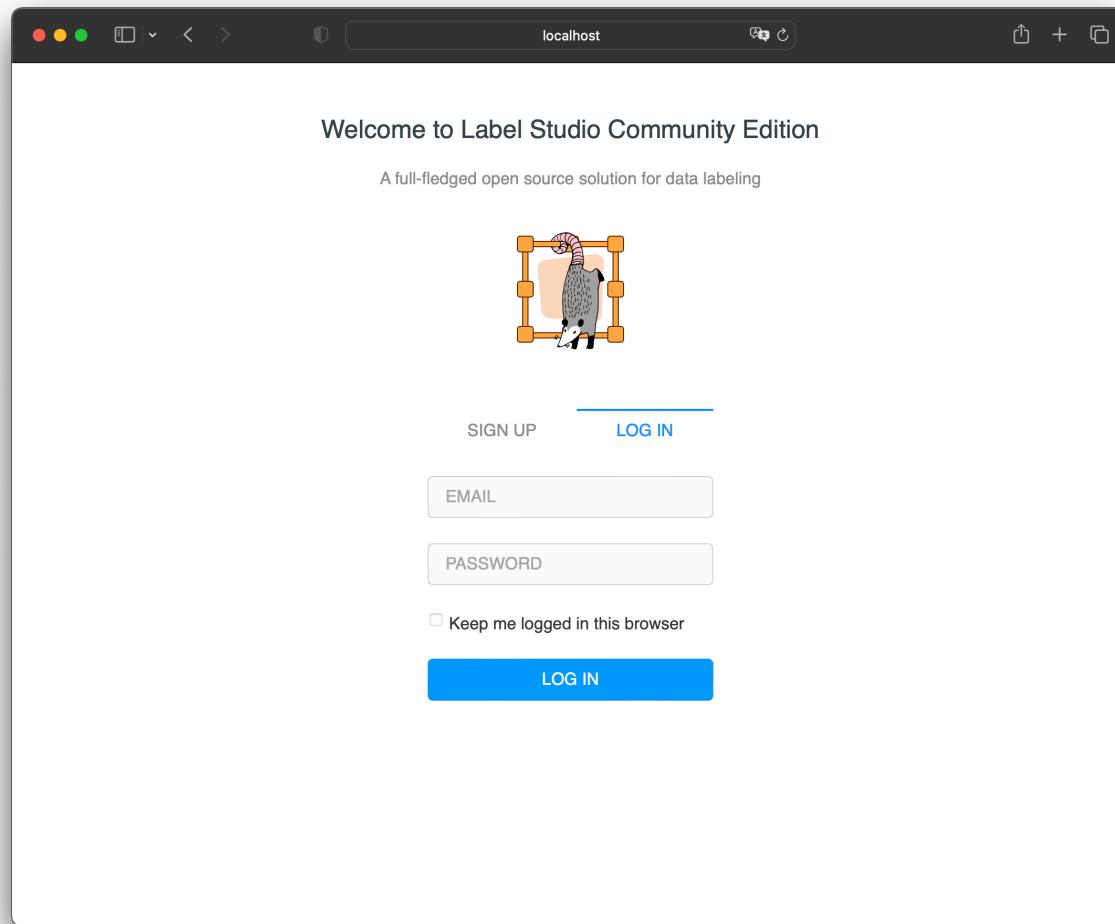
```
$ pip install label-studio
```

- Start Label Studio by entering the following command:

```
$ label-studio
=> Database and media directory: /Users/yilun/Library/Application Support/label-studio
=> Static URL is set to: /static/
...
Starting development server at http://0.0.0.0:8080/
Quit the server with CONTROL-C.
...
```

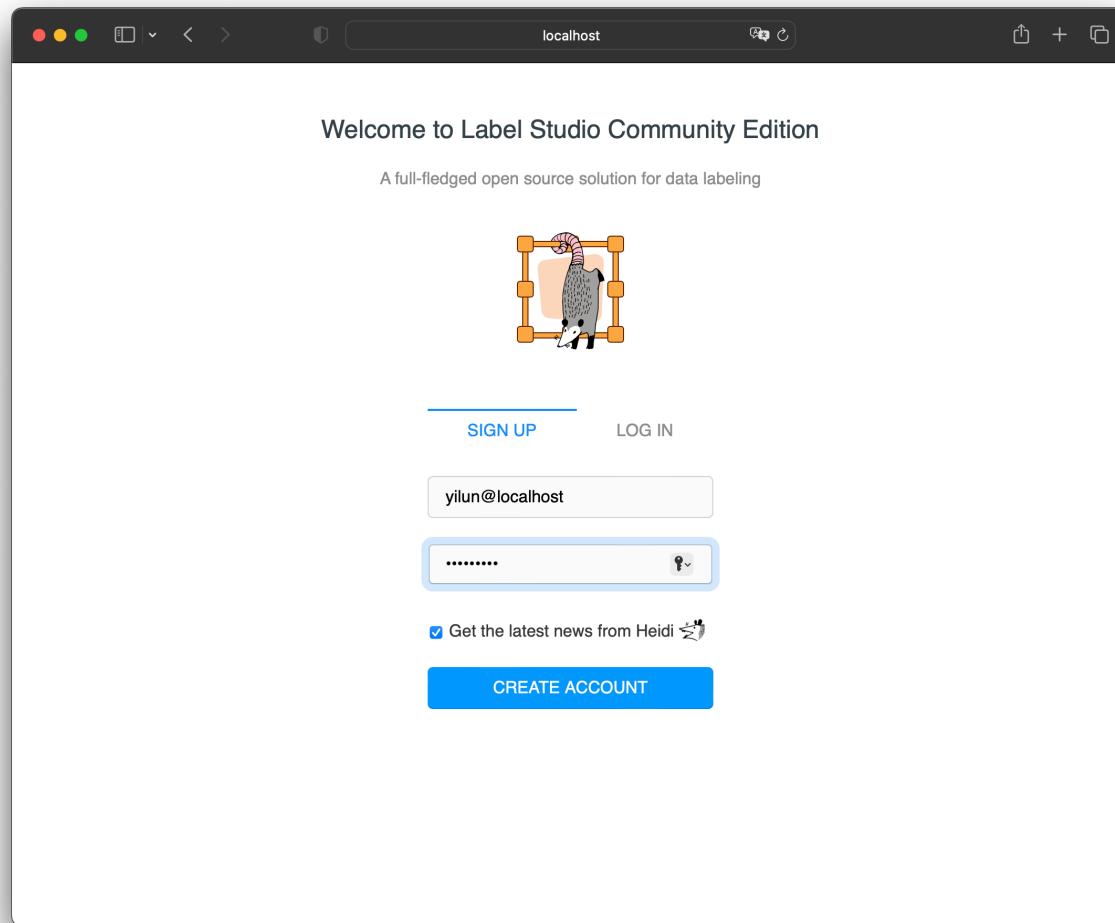
Step 2

- Visit localhost:8080 in your browser.



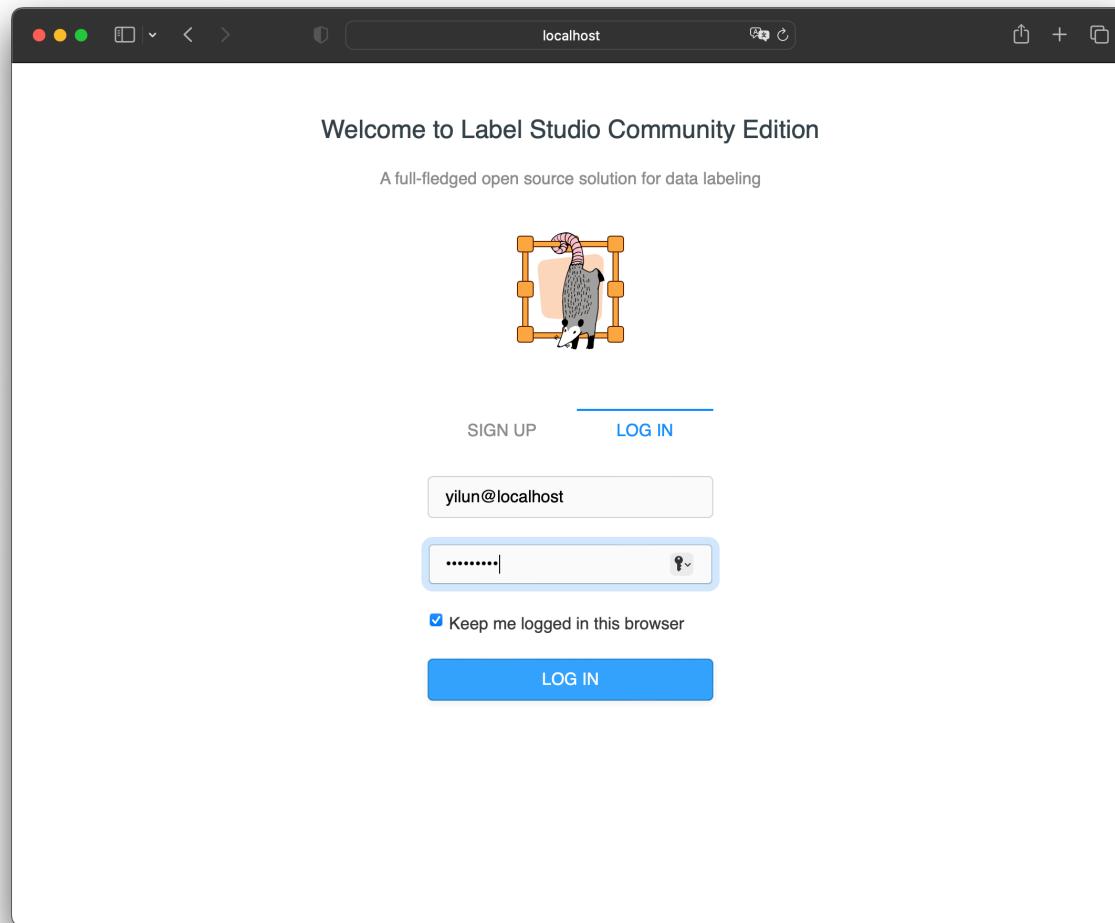
Step 3

- Sign up with an arbitrary email



Step 4

- Log in



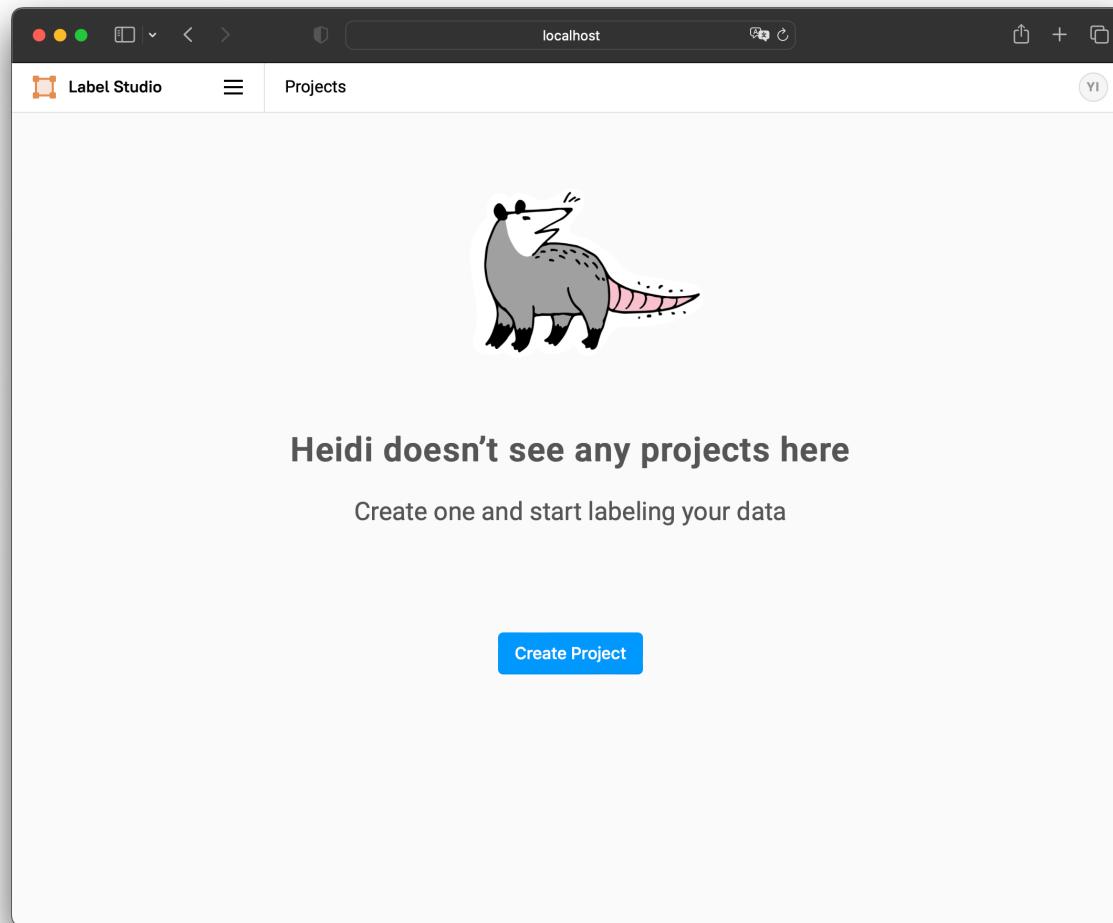
Step 5

- Collect 6 random dog and cat images from the internet.



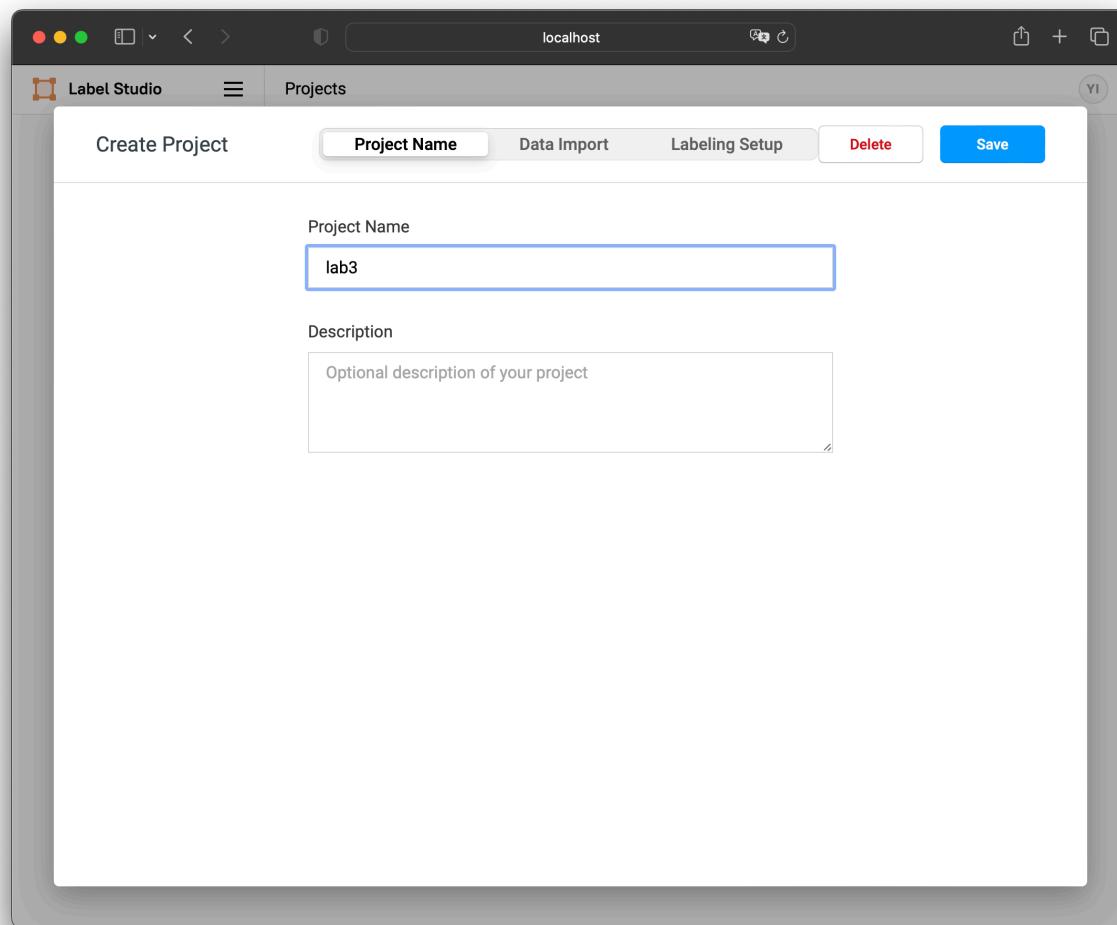
Step 6

- Create a project



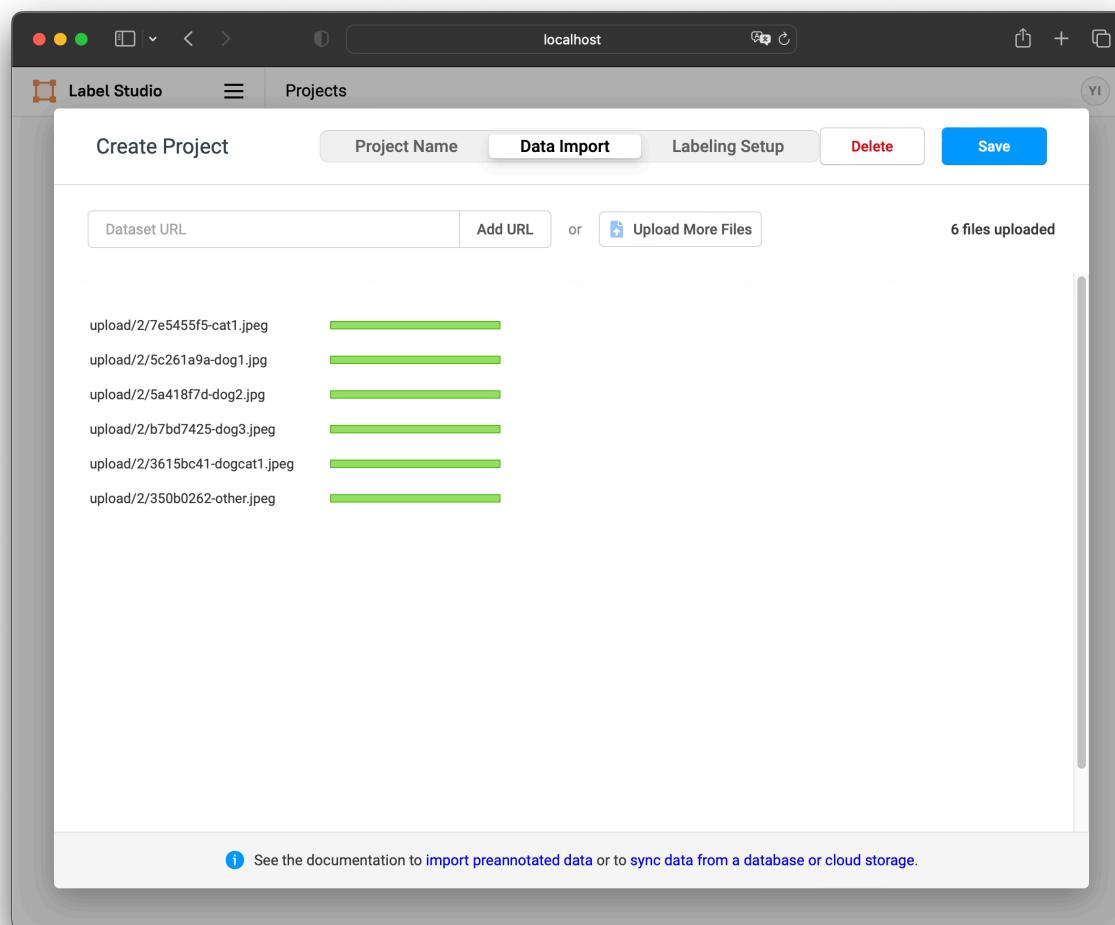
Step 7

- Set the project name



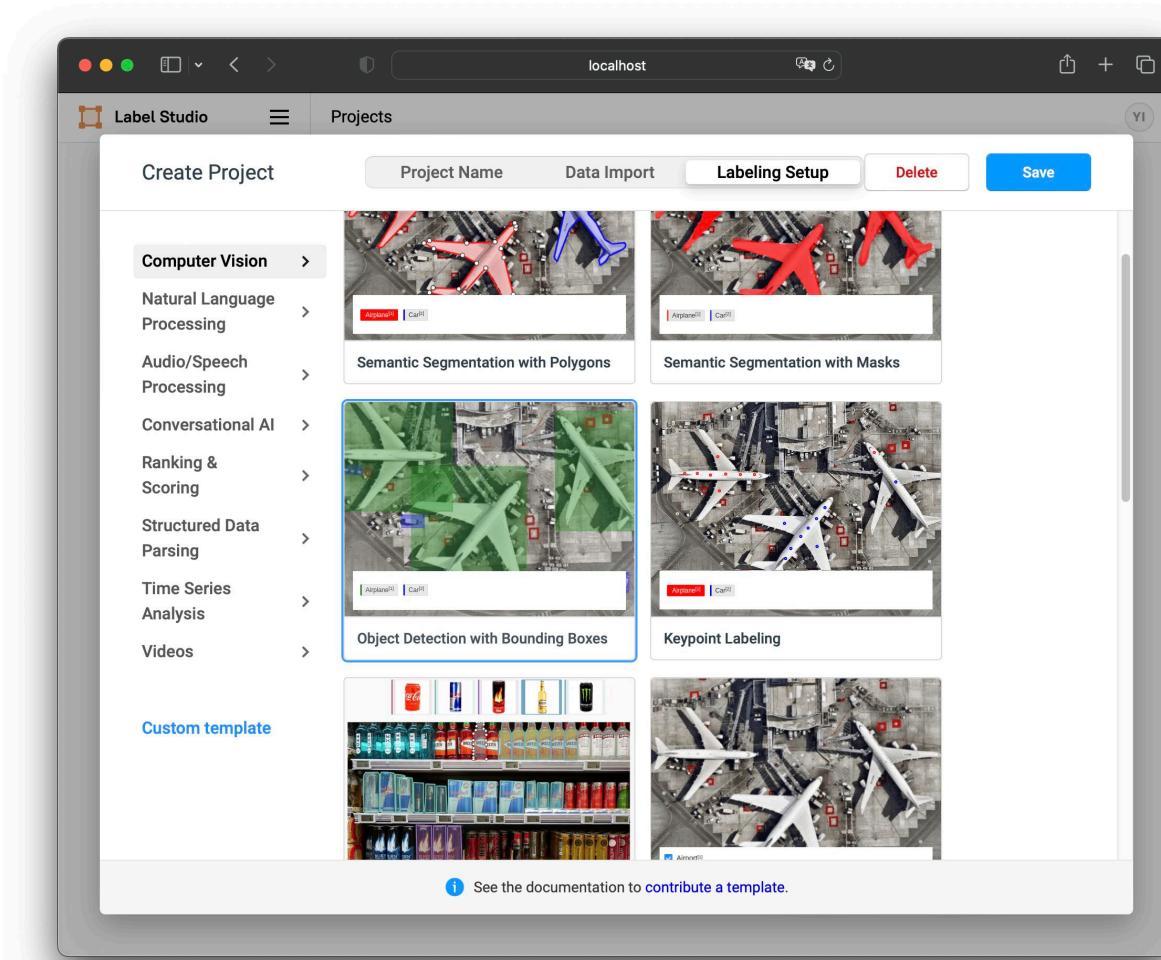
Step 8

- Use the Upload more files button to upload all dog and cat images.



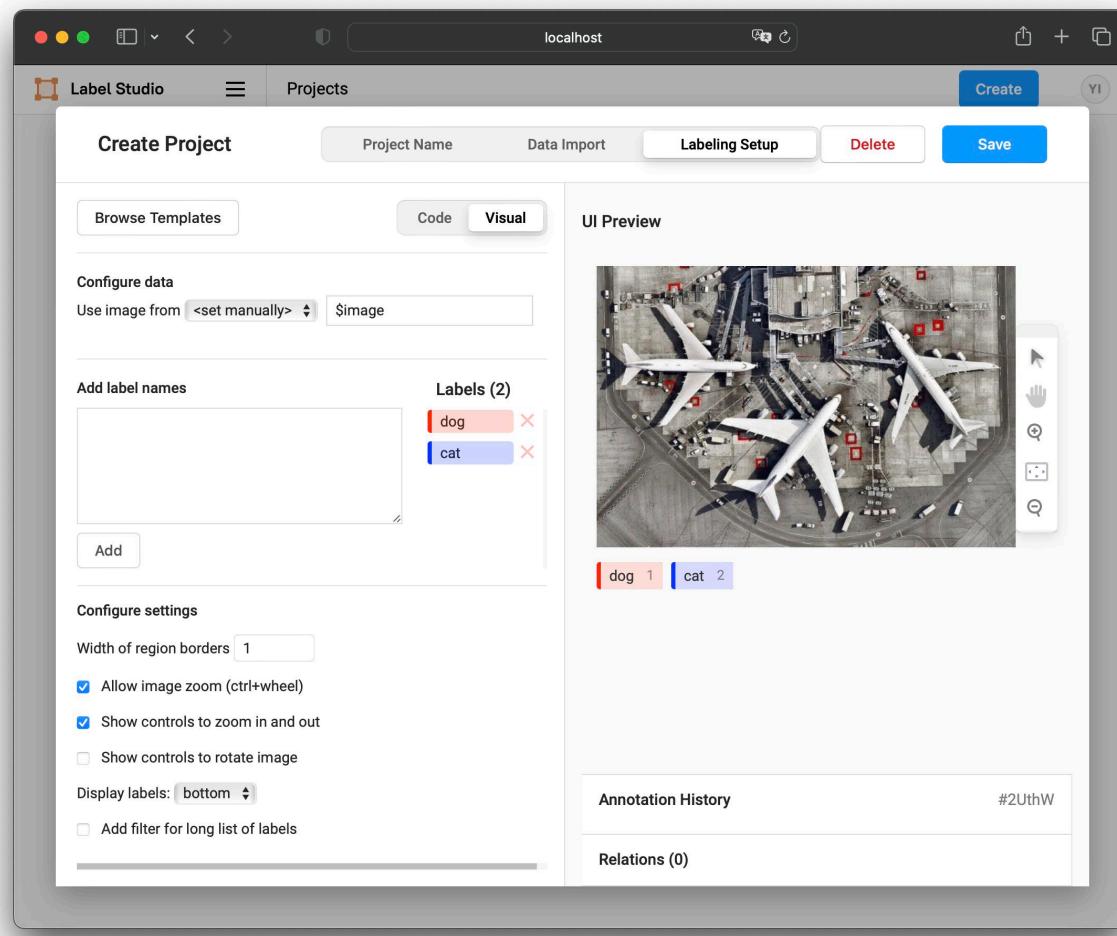
Step 9

- Select Object Detection with Bounding Boxes



Step 10

- Add dog and cat to label names. Note that the colors can be customized.
- Save.



Step 11

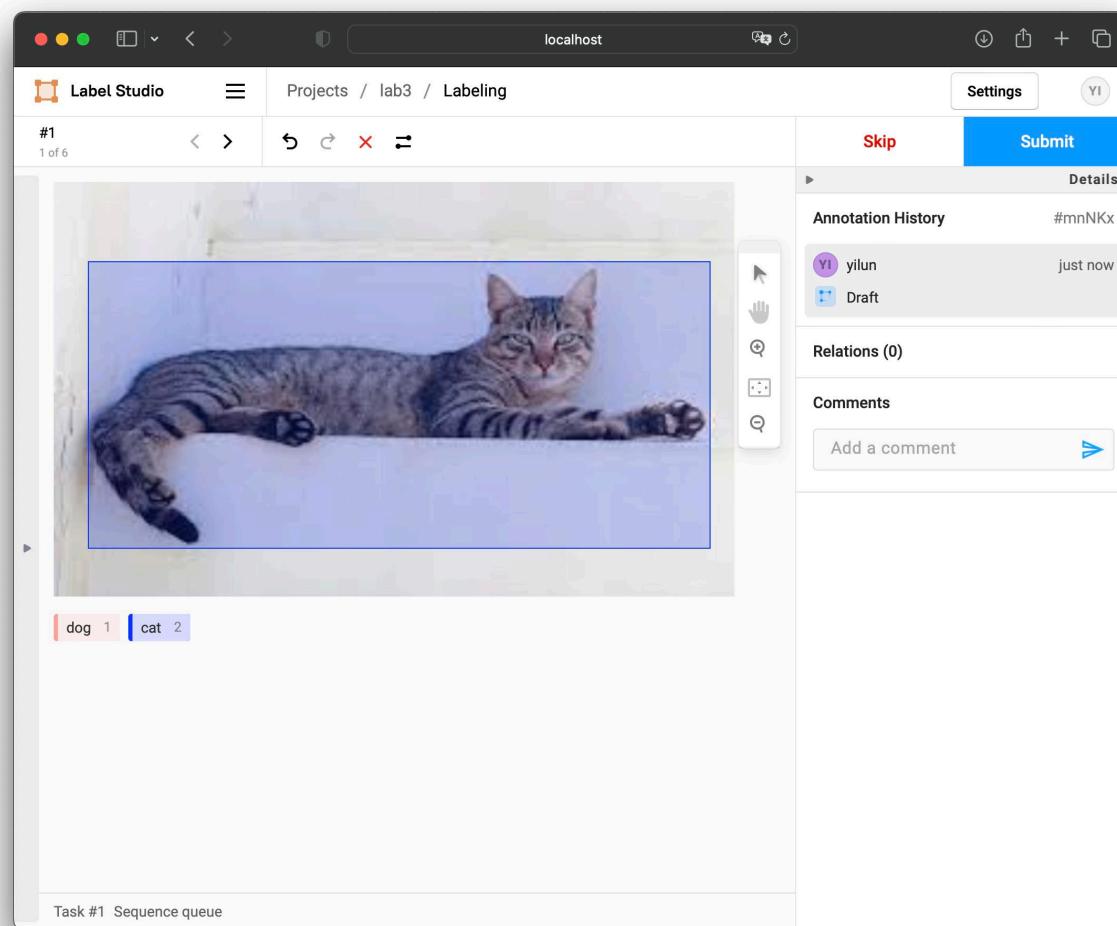
- Use the Label All Tasks button to start/resume labeling.

The screenshot shows the Label Studio interface running in a web browser on localhost. The project name is 'lab3'. The top navigation bar includes 'Actions', 'Columns', 'Filters', 'Order', 'not set', 'Label All Tasks' (which is highlighted in blue), 'Import', 'Export', 'List', and 'Grid'. Below this is a table with the following data:

ID	Completed	Annotated by	image
1	0	0	
2	0	0	
3	0	0	
4	0	0	
5	0	0	
6	0	0	

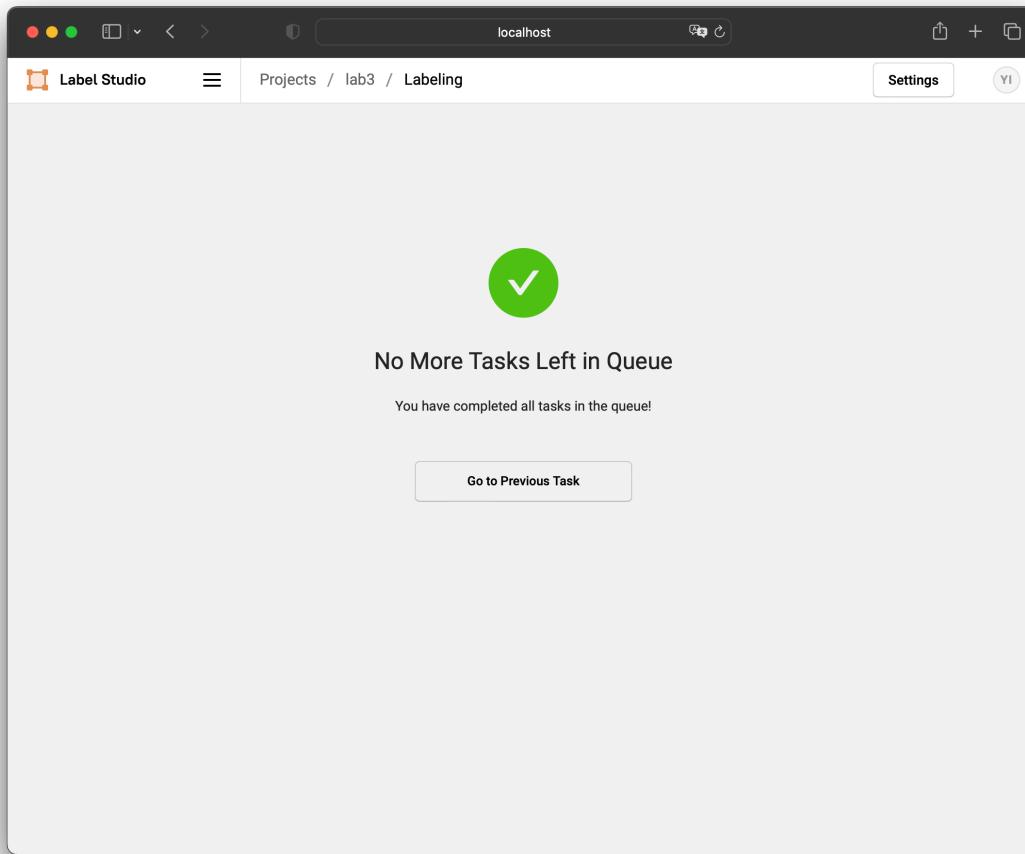
Step 12

- Select a label and drag the bounding box
- Using Submit to save the labels



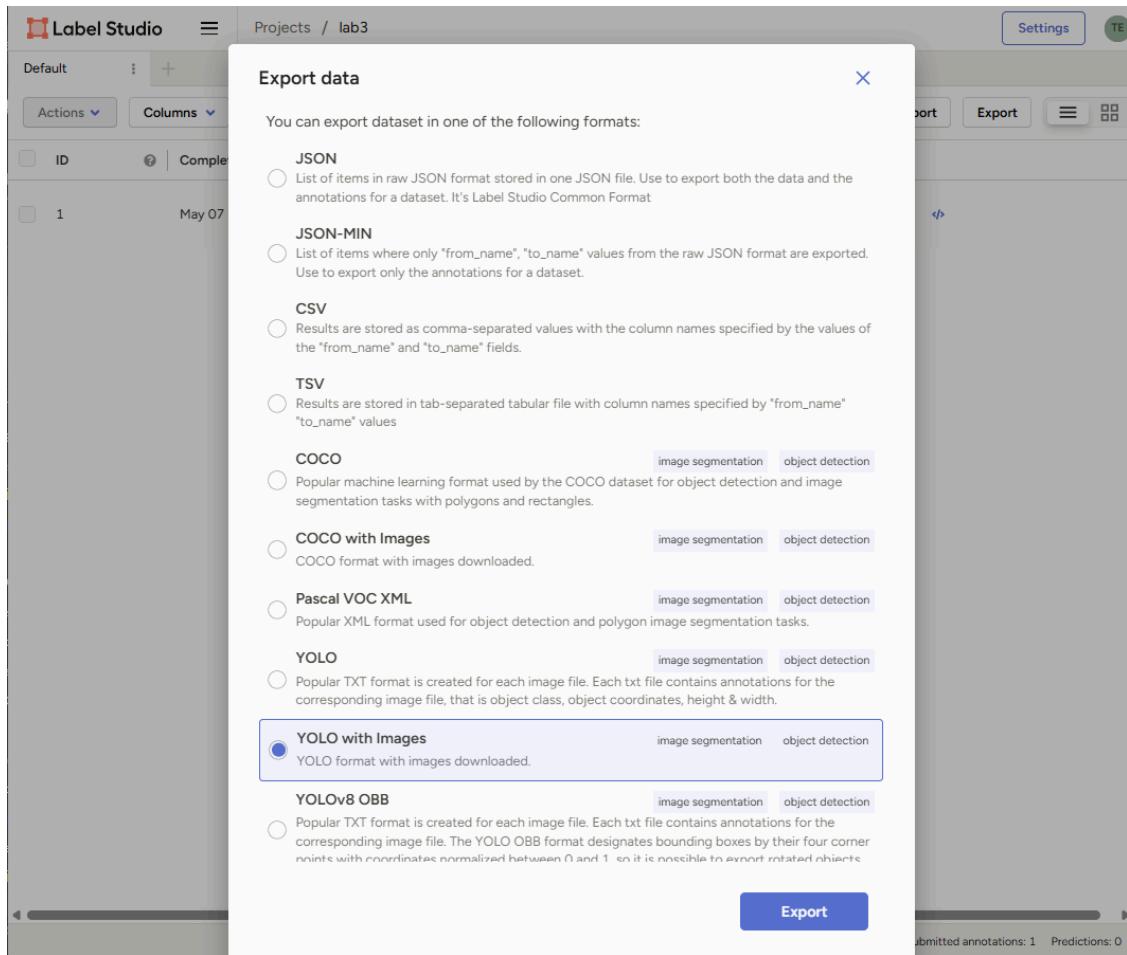
Step 13

- Finished!



Step 14

- Use the Export button to download labels and images.
- Select the **YOLO with Images** format, then export.



Transmit Dataset To Server

- On your computer, use `scp` to copy the data via the `ssh` protocol to the server.

```
$ scp -P 10122 \
  {path/to/file.zip} \
  {student_id}@140.113.160.149:~/hcc-ml/lab3
```

For example,

```
$ scp -P 10122 \
  ~/Downloads/project-4-at-2023-03-23-15-21-316bdd45.zip \
  123456@140.113.160.149:~/hcc-ml/lab3
```

Unzip the Dataset on the Server

- Unzip the dataset in the folder hcc-ml/lab-3 on the server:

```
$ unzip {path/to/file.zip} -d dogcat
```

- For example,

```
$ unzip project-4-at-2023-03-23-15-21-316bdd45.zip -d dogcat
```

Create a customized labeled file

- Create a YAML file called dogcat.yaml under the data folder:

```
# image root
train: ./dogcat/images
val: ./dogcat/images

# number of classes
nc: 2

# class names
names: ['cat', 'dog']
```

- **Important:** The order of names should follow the dogcat/classes.txt.

Create a customized model file

- Copy the file `cfg/training/yolov7-w6.yaml` to `cfg/training/dogcat-custom.yaml`

```
cp cfg/training/yolov7-w6.yaml cfg/training/dogcat-custom.yaml
```

- Modify the number of classes in `cfg/training/dogcat-custom.yaml`

```
nc: 80 -> nc: 2
```

Friendly Tips

- Use the tmux tool to start a new session window.

```
$ tmux new
```

- Run any long-running command in that session.
- Leave the session without interrupting the program by using the shortcut key:

```
push CTRL -> push B -> release CTRL and B -> push D -> release D
```

- Attach back to previously created session by using command

```
$ tmux attach
```

- Leave and delete the session by using the shortcut key:

```
push CTRL -> push D
```

- [Cheat sheet](#)

Start Training

- To start training, run the following command:

```
python train.py \
    --project dogcat \
    --epochs 50 \
    --device 0 \
    --batch-size 2 \
    --img 640 640 \
    --data data/dogcat.yaml \
    --cfg cfg/training/dogcat-custom.yaml \
    --weights weights/yolov7-w6.pt \
    --hyp data/hyp.scratch.custom.yaml
```

- Users can find weights and training info under dogcat/exp

Model Zoo

Model	Test Size	AP^{test}	$AP_{0.5}^{test}$	$AP_{0.75}^{test}$	batch 1 fps	batch 32 average time
YOLOv7-W6	1280	54.9%	72.6%	60.1%	84 <i>fps</i>	7.6 <i>ms</i>
YOLOv7-E6	1280	56.0%	73.5%	61.2%	56 <i>fps</i>	12.3 <i>ms</i>
YOLOv7-D6	1280	56.6%	74.0%	61.8%	44 <i>fps</i>	15.0 <i>ms</i>
YOLOv7-E6E	1280	56.8%	74.4%	62.1%	36 <i>fps</i>	18.7 <i>ms</i>

Checkpoint 3.

- After finishing the training, test the trained model with your images using the following command:

```
python demo.py \
--image dogcat/images/73770b2b-img4.jpg \
--weights dogcat/exp/weights/best.pt \
--conf_threshold 0.1
```

Replace 73770b2b-img4 with your own filename and select a proper threshold.

- Show demo.png which should contain some reasonable bounding boxes.



Friendly Tips Again

- Use nvtop to inspect the GPU and GPU memory usage of a server node.
- Use htop to inspect the CPU and system memory usage of a server node.
- When you log in, the login banner will show all GPU usage of this cluster (updated every 30 seconds).
- Each team can only use one GPU at a time.
- The user of GPU can be found in the login banner or using the command nvtop.
- If you find any team using more than one GPU, please report it to me (bkruan.ee11@nycu.edu.tw), and I will terminate their all process immediately.

The End